KAMRAN SHAIKH

kamran.aiml.engineer@gmail.com | sk2787434@gmail.com – +91 8591246016 – Mumbai, India

EXPERIENCE

Biz Technologies - (Decentrawood)
September 2025 – Present
AI/ML Engineer — Mumbai

- Built and deployed AI/ML solutions using the Gemini API for multimodal generation (image, video, music).
- Designed and optimized cryptocurrency price prediction models using time-series forecasting and deep learning (LSTM/GRU) to perform predictive market analysis.
- Developed NLP pipelines for text understanding and conversational AI use cases, including preprocessing, feature extraction, and model integration.
- Implemented and optimized Text-to-Speech (TTS) models for voice-enabled applications, focusing on speech quality and inference efficiency.
- Fine-tuned the DeepSeek R1 Large Language Model for domain-specific conversational and reasoning tasks, improving response accuracy and task alignment.
- Designed and developed scalable backend services (REST APIs) to serve AI/ML models, handling authentication, inference orchestration, and logging.
- Designed and developed scalable backend services (REST APIs) to serve AI/ML models, handling authentication, inference orchestration, and logging.
- Built and deployed AI/ML and Generative AI solutions on AWS, leveraging cloud services for scalability, monitoring, and production-grade reliability.

HealthIndia Insurance TPA Services Pvt. Ltd
April 2024 – September 2025
Machine Learning Engineer — Mumbai

- Developing AI/ML solutions for automating healthcare document processing, including claims forms, bills, and medical reports.
- Building end-to-end NLP and computer vision pipelines using YOLO, Azure OCR, and LLMs (LLaMA, GPT) for data extraction and classification.
- Designing scalable tools for automating NEFT extraction, CCN generation, and invoice digitization using Python, FastAPI, and OpenAI.
- Collaborating with cross-functional teams to deploy models and tools into production environments, improving operational efficiency and accuracy.

DG Market
Apr 2023 – Apr 2024
Python Developer — Mumbai

- Built Python-based automation tools and APIs for data extraction, transformation, and reporting.
- Developed robust pipelines for scraping, processing, and storing large datasets used in global tender analytics.

Ardentisys-Tebillion
October 2022 – March 2023
Associate Engineer — Mumbai

- Contributed to the development of business automation solutions, focusing on backend logic, API integrations, and process optimization using Python and automation tools.

- Supported product teams in building scalable software modules for client-driven enterprise solutions.

SKILLS

- LLMs & AI Agents: Fine-tuning, Retrieval-Augmented Generation (RAG), Tool-using Agents, Chatbots, LLM Prompts, GPT-4, LLaMA, Gemma, Bloom, Ollama, LangChain, LangGraph,ChromaDB, FAISS
- Machine Learning & Deep Learning: Supervised/unsupervised ML, neural networks, CNNs, LSTMs, GANs,Decision Trees,Random Forest, Gradient Boosting (XGBoost)„Support Vector Regression (SVR), Libraries: Scikit-learn, TensorFlow, PyTorch, Keras
- NLP & Document AI: Named Entity Recognition (NER), Text Classification, OCR, Document Parsing, Tools: SpaCy, BERT, Azure OCR, Tesseract, OpenCV, YOLO, Stable Diffusion
- Programming & Automation: Languages: Python , SQL, Tools: Pandas, NumPy, Selenium, Regex, FastAPI
- Deployment & Infrastructure: Experience with cloud and on-prem environments, AWS EC2 & S3, Linux servers, MongoDB
- Data Engineering: Web scraping, data pipelines, SQL-based data storage and integration, SQL, ER diagrams, data validation checks

PROJECTS

Document Classification System YOLO, CNN, LSTM, GRU, Python
Automated classification of insurance claim documents using a multi-stage pipeline combining object detection, image classification, and text classification.

Auto NEFT Extraction YOLOv8s, Tesseract OCR, Python
Extracted NEFT data from scanned cheques using object detection, OCR, and regex-based postprocessing for structured output.

Auto CCN Generation YOLO, Azure OCR, LLM (LLaMA/Gemma), Python
Developed an AI pipeline to extract data from healthcare claim forms and generate Control Claim Numbers (CCNs) automatically.

Auto Bill Entry System LLaMA 3.1, Python
Built a system to extract structured data from invoices and generate JSON outputs for seamless database integration.

Aadhaar Card Masking Tool YOLOv8s, OpenCV, Python
Designed a privacy tool to detect and redact Aadhaar numbers from scanned documents using real-time object detection.

Medical NER System SpaCy, BERT, Python
Trained custom NER models to extract patient names, diagnoses, treatments, and doctor information from medical records.

ITVx – Web Scraper & Data Pipeline Selenium, Python, AWS S3, SQL
Built an automated web scraper and data pipeline to extract, clean, and store structured data.

Chatbot with Memory GPT-API, LangChain, OpenAI, Python
Created an intelligent chatbot using LangChain agents capable of memory, context retention, and tool usage.

PDF Chatbot LangChain, ChromaDB, OpenAI, PyMuPDF
Built a document Q&A chatbot that can ingest and answer questions from PDF files.

Sentiment Analysis CountVectorizer, Scikit-learn, Python
Implemented a machine learning model to classify customer sentiments in product reviews.

Face Recognition Attendance System OpenCV, Dlib, Python
A real-time facial recognition system for automated employee attendance tracking.

Stock Market Prediction Python,, Pandas, Matplotlib, LSTM
Developed a predictive analytics platform combining deep learning and machine learning models to analyze cryptocurrency market trends and visualize forecasts and building real-time trading bots

Multimodal Generative AI Platform Gemini API, Python, FastAPI, AWS, MongoDB
Built an end-to-end multimodal AI platform enabling text-to-image, image-to-video, and music generation using the Gemini API, with backend APIs for request handling, metadata storage, and inference orchestration.

LLM Fine-Tuning & Deployment System (DeepSeek R1) DeepSeek R1, PyTorch, Hugging Face, AWS, MongoDB
Worked with cloud-based GPUs (NVIDIA A100, T4, RTX 4090) to train, fine-tune the DeepSeek R1 Large Language Model on domain-specific datasets and deployed it as a scalable inference service, enabling accurate reasoning and conversational AI for enterprise use cases.

AI Inference Backend & Model Serving Platform FastAPI, Docker, AWS ECS, MongoDB
Designed and implemented a backend system for serving multiple AI models (LLMs, CV, NLP), including request routing, versioning, logging, and performance monitoring.

EDUCATION

M.H. Saboo Siddik College of Engineering 2015 – 2021
B.E. in Electronics Engineering Mumbai

Thakur Polytechnic College of Engineering 2012 – 2015
Diploma in Electronics Engineering Mumbai

CERTIFICATIONS

Master's Certification in Data Science
Covered machine learning, deep learning, NLP, Python, and real-world data science projects.